

# Socket Programming

# CS 353 – Computer Networks Lab

1. Labs 1 and 2 are not graded.
2. All the rest of labs are graded.
3. One week time to evaluation.
4. NetSim software is used.
5. Grading: 50% lab sessions.

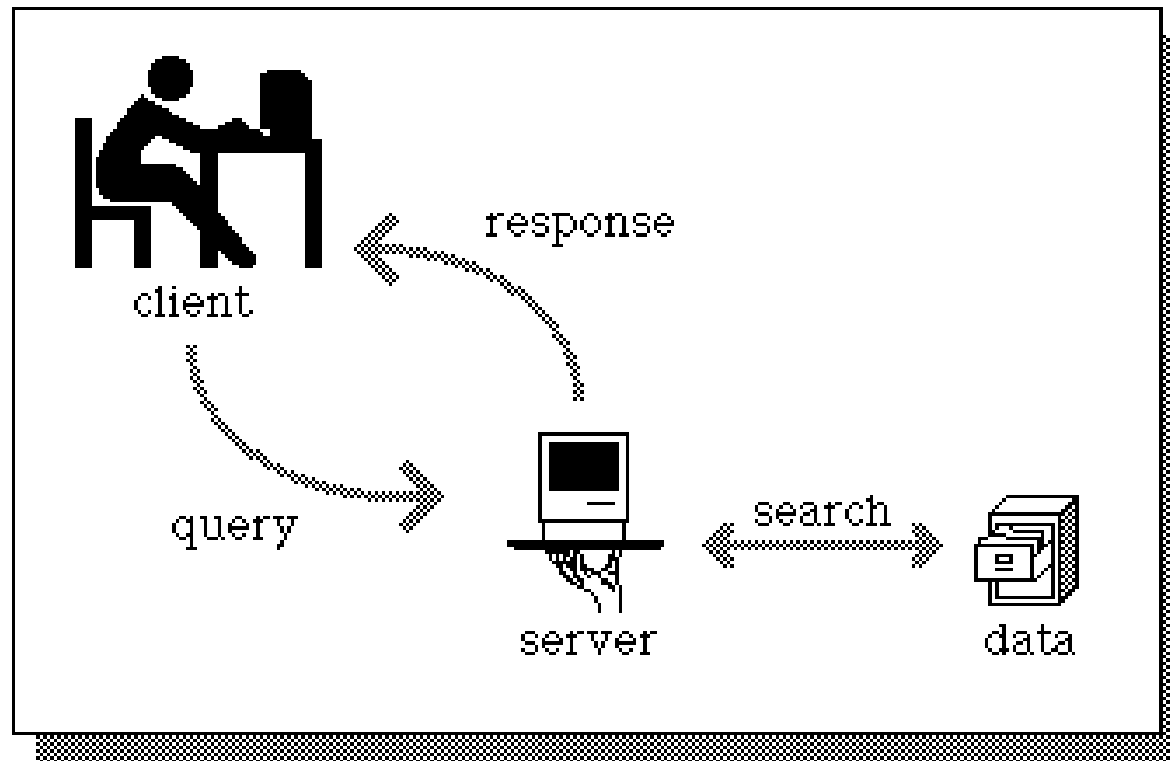
30% end-sem

20% lab-quiz

# Client-Server Architecture

- In the client server architecture, a machine(referred as client) makes a request to connect to another machine (called as server) for providing some service.
- The services running on the server **run on known ports**(application identifiers) and the client needs to know the address of the server machine and this port in order to connect to the server.
- On the other hand, the server does not need to know about the address or the port of the client at the time of connection initiation.
- The first packet which the client sends as a request to the server contains these informations about the client which are further used by the server to send any information.
- Client acts as the active device which makes the first move to establish the connection whereas the server passively waits for such requests from some client.

# Client-Server Architecture



# What is a Socket?

- In unix, whenever there is a need for inter process communication within the same machine, we use mechanism like signals or pipes.
- Similarly, when we desire a communication between two applications possibly running on different machines, we need sockets.
- Sockets are treated as another entry in the unix open file table. So all the system calls which can be used for any IO in unix can be used on socket.
- The server and client applications use various system calls to connect which use the basic construct called socket.
- A socket is one end of the communication channel between two applications running on different machines.

- Steps followed by client to establish the connection:

- 1) Create a socket
- 2) Connect the socket to the address of the server
- 3) Send/Receive data
- 4) Close the socket

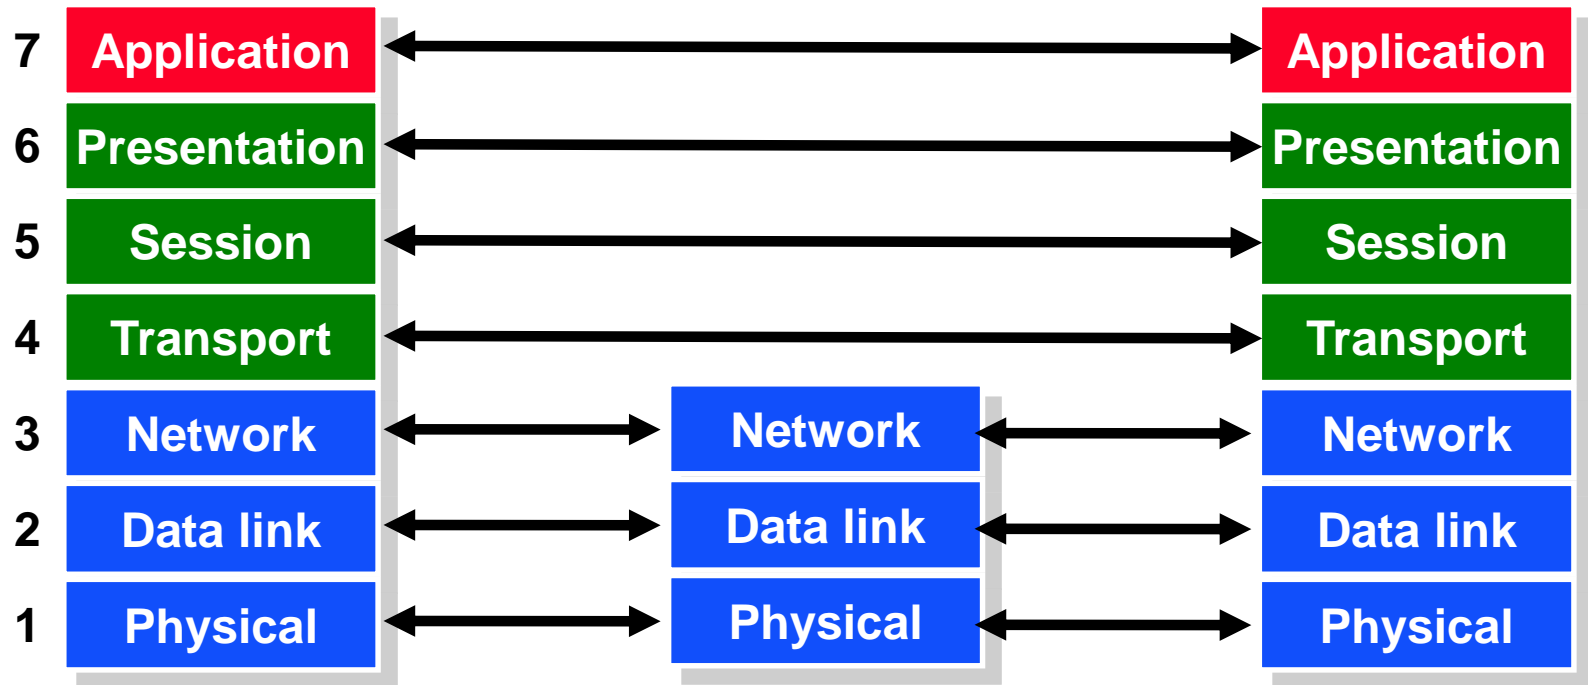
- Steps followed by server to establish the connection:

- 1) Create a socket
- 2) Bind the socket to the port number known to all clients
- 3) Listen for the connection request
- 4) Accept connection request
- 5) Send/Receive data

# Why Socket?

- How can I program a network application?
  - Share data
  - Send messages
  - Finish course projects...
- IPC - Interprocess Communication

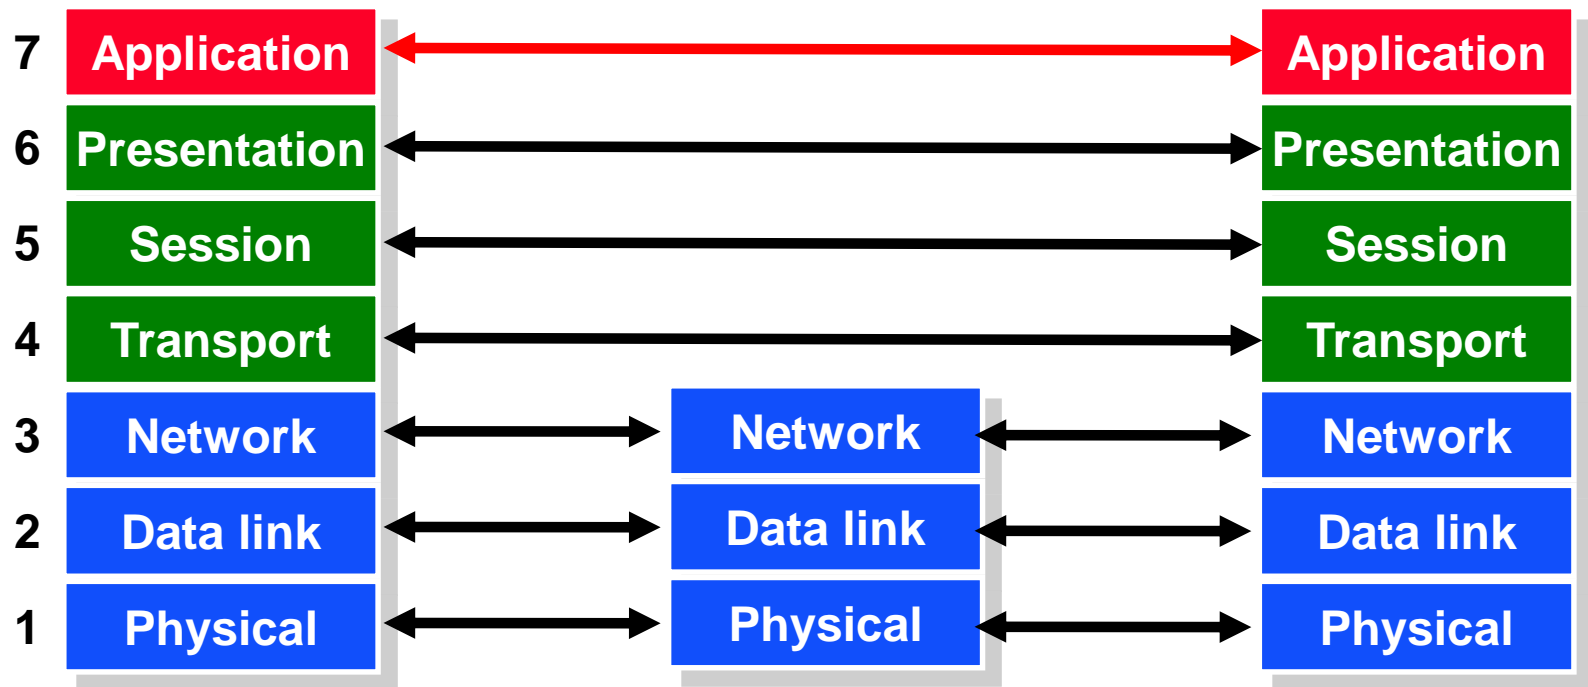
# Network Layering





# Network Layering

- Why layering?



# Layering Makes it Easier

- Application programmer
  - Doesn't need to send IP packets
  - Doesn't need to send Ethernet frames
  - Doesn't need to know how TCP implements reliability
- Only need a way to pass the data down
  - Socket is the API to access transport layer functions

# What Lower Layer Need to Know?


- We pass the data down. What else does the lower layer need to know?

# What Lower Layer Need to Know?

- We pass the data down. What else does the lower layer need to know?
- How to identify the destination process?
  - Where to send the data? (Addressing)
  - What process gets the data when it is there? (Multiplexing)

# Identify the Destination

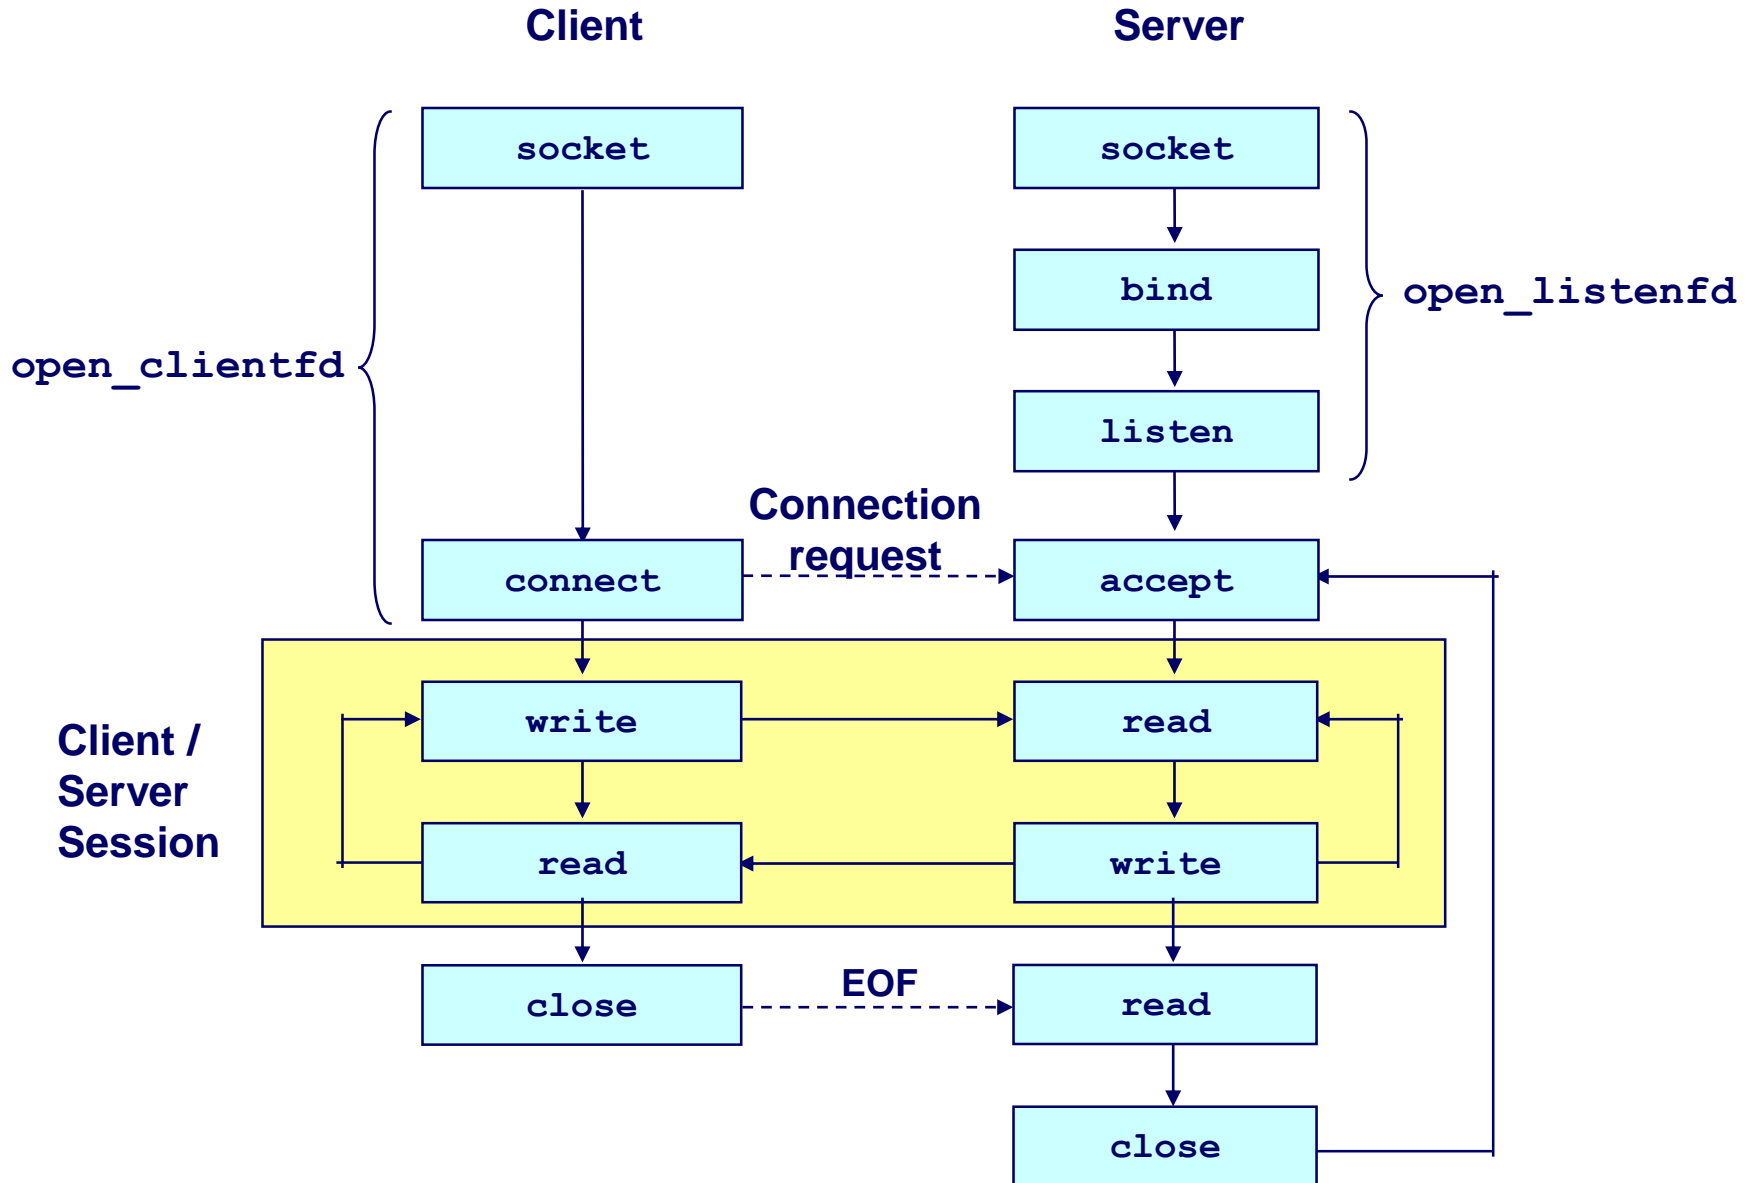
- Addressing
  - IP address
  - hostname (resolve to IP address via DNS)
- Multiplexing
  - port



# Sockets

- How to use sockets
  - Setup socket
    - Where is the remote machine (IP address, hostname)
    - What service gets the data (port)
  - Send and Receive
    - Designed just like any other I/O in unix
    - send -- write
    - recv -- read
  - Close the socket

# Overview



# Basic data structures used in Socket programming

- Socket Descriptor: A simple file descriptor in Unix.  
Int
- Socket Address: This construct holds the information for socket address

```
struct sockaddr {  
    unsigned short  sa_family;  // address family, AF_XXX  
    or PF_XXX  
    char  sa_data[14]; // 14 bytes of protocol address  
};
```

- AF stands for Address Family and PF stands for Protocol Family. In most modern implementations only the AF is being used.

Name	Purpose
AF_UNIX, AF_LOCAL	Local communication
AF_INET	IPv4 Internet protocols
AF_INET6	IPv6 Internet protocols



# Basic data structures used in Socket programming

- Socket Descriptor: A simple file descriptor in Unix.  
Int
- Socket Address: This construct holds the information for socket address

```
struct sockaddr {  
    unsigned short  sa_family;  // address family, AF_xxx  
    or PF_xxx  
    char  sa_data[14]; // 14 bytes of protocol address  
};
```

- AF stands for Address Family and PF stands for Protocol Family. In most modern implementations only the AF is being used.

Name	Purpose
AF_UNIX, AF_LOCAL	Local communication
AF_INET	IPv4 Internet protocols
AF_INET6	IPv6 Internet protocols

# Step 1 – Setup Socket

- **Both client and server need to setup the socket**
  - *int socket(int domain, int type, int protocol);*
- *domain*
  - AF\_INET -- IPv4 (AF\_INET6 for IPv6)
- *type*
  - SOCK\_STREAM -- TCP
  - SOCK\_DGRAM -- UDP
- *protocol*
  - 0
- For example,
  - *int sockfd = socket(AF\_INET, SOCK\_STREAM, 0);*

# Step 2 (Server) - Binding

- **Only server need to bind**
  - *int bind(int sockfd, const struct sockaddr \*my\_addr, socklen\_t addrlen);*
- *sockfd*
  - file descriptor socket() returned
- *my\_addr*
  - struct sockaddr\_in for IPv4
  - cast (struct sockaddr\_in\*) to (struct sockaddr\*)

```
struct sockaddr_in {
    short    sin_family; // e.g. AF_INET
    unsigned short sin_port; // e.g. htons(3490)
    struct in_addr sin_addr; // see struct in_addr, below
    char      sin_zero[8]; // zero this if you want to
};
struct in_addr {
    unsigned long s_addr; // load with inet_aton()
};
```

# What is that Cast?

- bind() takes in protocol-independent (struct sockaddr\*)

```
struct sockaddr {  
    unsigned short  sa_family; // address family  
    char            sa_data[14]; // protocol address  
};
```

— There are structs for IPv6, etc.

# IP Addresses

Assuming that we are dealing with IPv4 addresses, the address is a 32bit integer. Remembering a 32 bit number is not convenient for humans. So, the address is written as a set of four integers seperated by dots, where each integer is a representation of 8 bits.

The representation is like a.b.c.d, where a is the representation of the most significant byte. The system call which converts this representation into Network Byte Order is:

**int inet\_aton(const char \*cp, struct in\_addr \*inp);**

inet\_aton() converts the Internet host address cp from the standard numbers-and-dots notation into binary data and stores it in the structure that inp points to. inet\_aton returns nonzero if the address is valid, zero if not.

For example, if we want to initialize the sockaddr\_in construct by the IP address and desired port number, it is done as follows:

```
struct sockaddr_in sockaddr;  
sockaddr.sin_family = AF_INET;  
sockaddr.sin_port = htons(21);  
inet_aton("172.26.117.168", &(sockaddr.sin_addr));  
memset(&(sockaddr.sin_zero), '\0', 8);
```

# Step 2 (Server) - Binding contd.

- *addrlen*
  - size of the `sockaddr_in`

```
struct sockaddr_in saddr;  
int sockfd;  
unsigned short port = 80;
```

```
if((sockfd=socket(AF_INET, SOCK_STREAM, 0) < 0) {      // from back a couple slides  
    printf("Error creating socket\n");  
    ...  
}
```

```
memset(&saddr, '\0', sizeof(saddr));                // zero structure out  
saddr.sin_family = AF_INET;                          // match the socket() call  
saddr.sin_addr.s_addr = htonl(INADDR_ANY);           // bind to any local address  
saddr.sin_port = htons(port);                        // specify port to listen on
```

```
if((bind(sockfd, (struct sockaddr *) &saddr, sizeof(saddr)) < 0) { // bind!  
    printf("Error binding\n");  
    ...  
}
```

# What is htonl(), htons()?

- Byte ordering
  - Network order is big-endian
  - Host order can be big- or little-endian
    - x86 is little-endian
    - SPARC is big-endian
- Conversion
  - *htons()*, *htonl()*: host to network short/long
  - *ntohs()*, *ntohl()*: network order to host short/long
- What need to be converted?
  - Addresses
  - Port
  - etc.

# Little & Big Endian

- Some systems (like x8086) are Little Endian i-e. least significant byte is stored in the higher address, whereas in Big endian systems most significant byte is stored in the higher address.
- Consider a situation where a Little Endian system wants to communicate with a Big Endian one, if there is no standard for data representation then the data sent by one machine is misinterpreted by the other.
- So standard has been defined for the data representation in the network (called Network Byte Order) which is the Big Endian.
- The system calls that help us to convert a short/long from Host Byte order to Network Byte Order and vice-versa are
  - 1) htons() -- "Host to Network Short"
  - 2) htonl() -- "Host to Network Long"
  - 3) ntohs() -- "Network to Host Short"
  - 4) ntohl() -- "Network to Host Long"



# Step 3 (Server) - Listen

- **Now we can listen**
  - *int listen(int sockfd, int backlog);*
- *sockfd*
  - again, file descriptor socket() returned
- *backlog*
  - number of pending connections to queue
- For example,
  - *listen(sockfd, 5);*

# Step 4 (Server) - Accept

- **Server must explicitly accept incoming connections**
  - *int accept(int sockfd, struct sockaddr \*addr, socklen\_t \*addrlen)*
- *sockfd*
  - again... file descriptor socket() returned
- *addr*
  - pointer to store client address, (struct sockaddr\_in \*) cast to (struct sockaddr \*)
- *addrlen*
  - pointer to store the returned size of addr, should be sizeof(\*addr)
- For example
  - *int isock=accept(sockfd, (struct sockaddr\_in \*) &caddr, &crlen);*

# Put Server Together

```
struct sockaddr_in saddr, caddr;
int sockfd, clen, isock;
unsigned short port = 80;

if((sockfd=socket(AF_INET, SOCK_STREAM, 0) < 0) { // from back a couple slides
    printf("Error creating socket\n");
    ...
}

memset(&saddr, '\0', sizeof(saddr)); // zero structure out
saddr.sin_family = AF_INET; // match the socket() call
saddr.sin_addr.s_addr = htonl(INADDR_ANY); // bind to any local address
saddr.sin_port = htons(port); // specify port to listen on

if((bind(sockfd, (struct sockaddr *) &saddr, sizeof(saddr)) < 0) { // bind!
    printf("Error binding\n");
    ...
}

if(listen(sockfd, 5) < 0) { // listen for incoming connections
    printf("Error listening\n");
    ...
}

clen=sizeof(caddr)
if((isock=accept(sockfd, (struct sockaddr *) &caddr, &clen)) < 0) { // accept one
    printf("Error accepting\n");
    ...
}
```

# What about client?

- Client need not bind, listen, and accept
- **All client need to do is to connect**
  - *int connect(int sockfd, const struct sockaddr \*saddr, socklen\_t addrlen);*
- For example,
  - *connect(sockfd, (struct sockaddr \*) &saddr, sizeof(saddr));*

# Domain Name System (DNS)

- What if I want to send data to “www.slashdot.org”?
  - DNS: Conceptually, DNS is a database collection of host entries

```
struct hostent {  
    char      *h_name;    // official hostname  
    char      **h_aliases; // vector of alternative hostnames  
    int  h_addrtype;    // address type, e.g. AF_INET  
    int  h_length;      // length of address in bytes, e.g. 4 for IPv4  
    char      **h_addr_list; // vector of addresses  
    char      *h_addr;    // first host address, synonym for h_addr_list[0]  
};
```

- hostname -> IP address
  - *struct hostent \*gethostbyname(const char \*name);*
- IP address -> hostname
  - *struct hostent \*gethostbyaddr(const char \*addr, int len, int type);*

# Put Client Together

```
struct sockaddr_in saddr;
struct hostent *h;
int sockfd, connfd;
unsigned short port = 80;

if((sockfd=socket(AF_INET, SOCK_STREAM, 0) < 0) { // from back a couple slides
    printf("Error creating socket\n");
    ...
}

if((h=gethostbyname("www.slashdot.org")) == NULL) { // Lookup the hostname
    printf("Unknown host\n");
    ...
}

memset(&saddr, '\0', sizeof(saddr)); // zero structure out
saddr.sin_family = AF_INET; // match the socket() call
memcpy((char *) &saddr.sin_addr.s_addr, h->h_addr_list[0], h->h_length); // copy the address
saddr.sin_port = htons(port); // specify port to connect to

if((connfd=connect(sockfd, (struct sockaddr *) &saddr, sizeof(saddr)) < 0) { // connect!
    printf("Cannot connect\n");
    ...
}
```

# We Are Connected

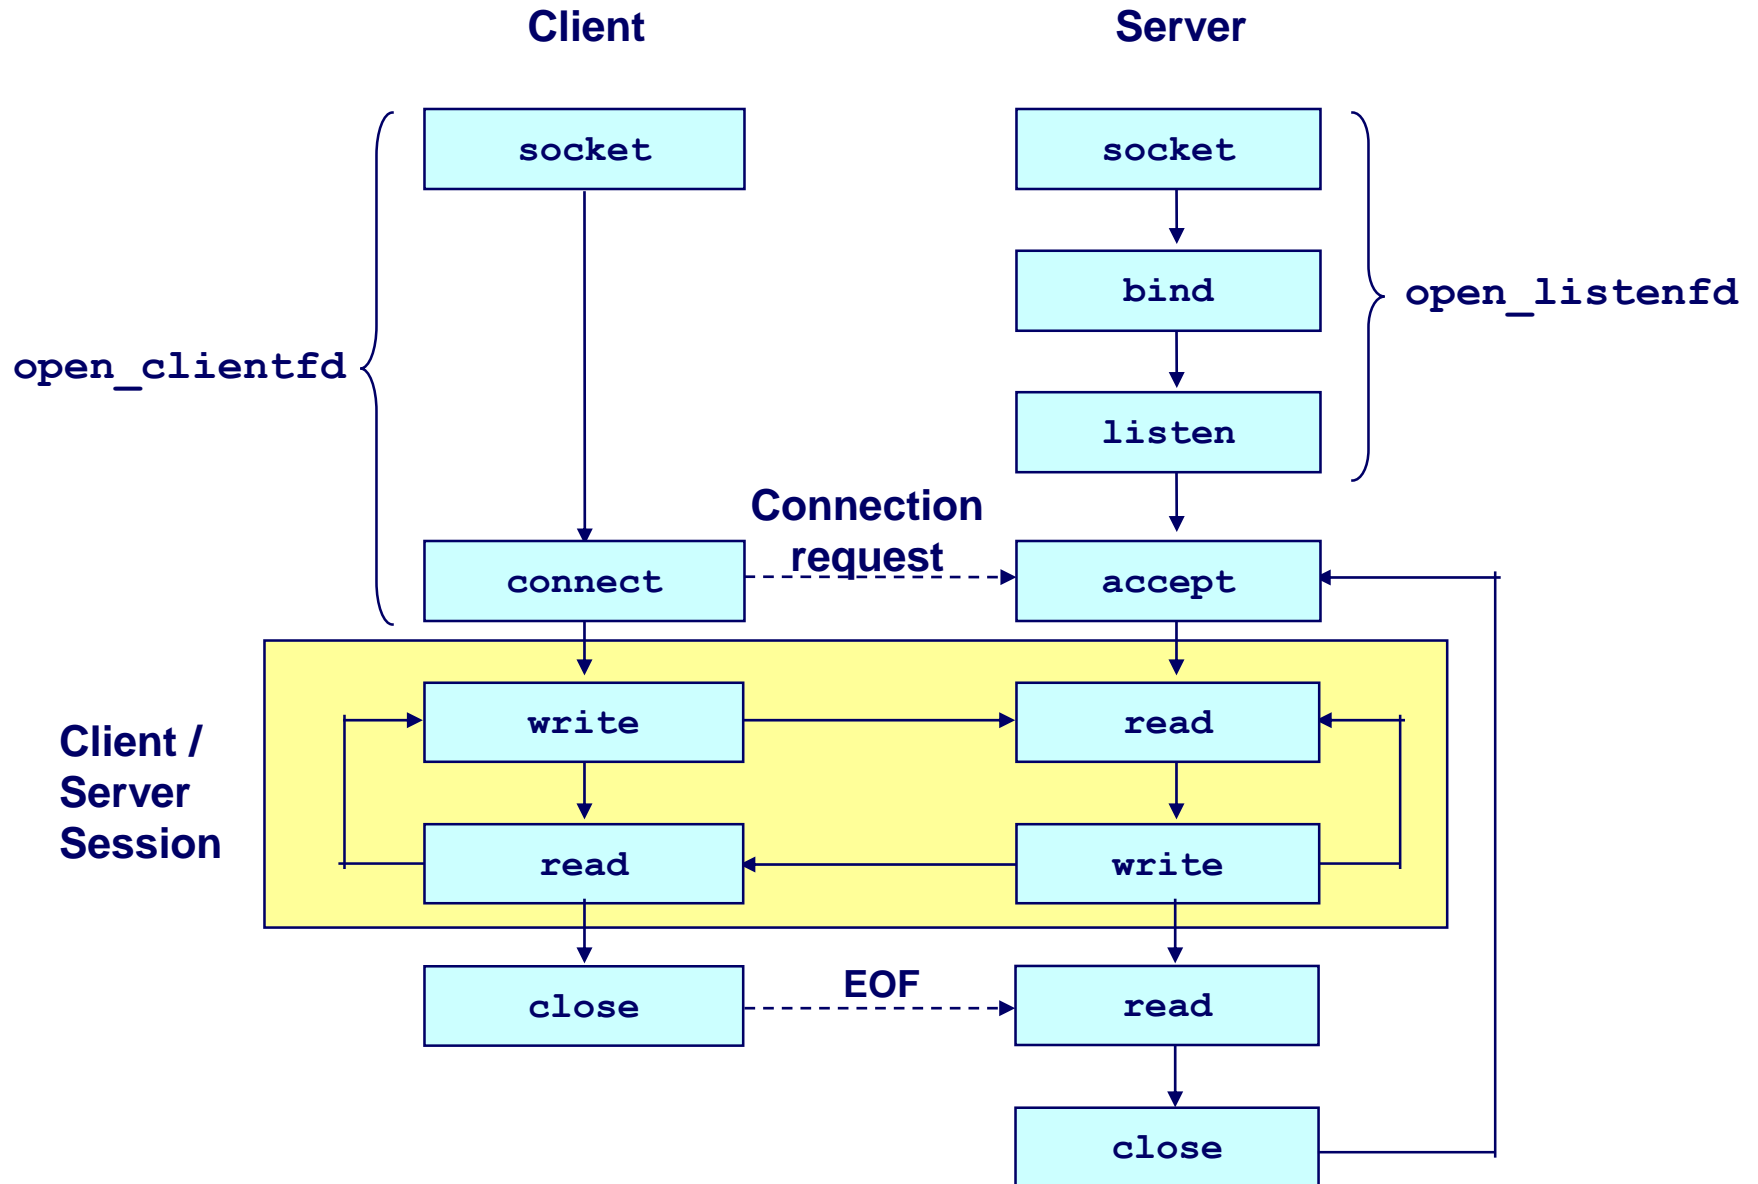
- Server accepting connections and client connecting to servers
- Send and receive data
  - *ssize\_t read(int fd, void \*buf, size\_t len);*
  - *ssize\_t write(int fd, const void \*buf, size\_t len);*
- For example,
  - *read(sockfd, buffer, sizeof(buffer));*
  - *write(sockfd, "hey\n", strlen("hey\n"));*

# TCP Framing

- TCP does NOT guarantee message boundaries
  - IRC commands are terminated by a newline
  - But you may not get one at the end of read(), e.g.
    - One Send “Hello\n”
    - Multiple Receives “He”, “llo\n”
  - If you don’t get the entire line from one read(), use a buffer



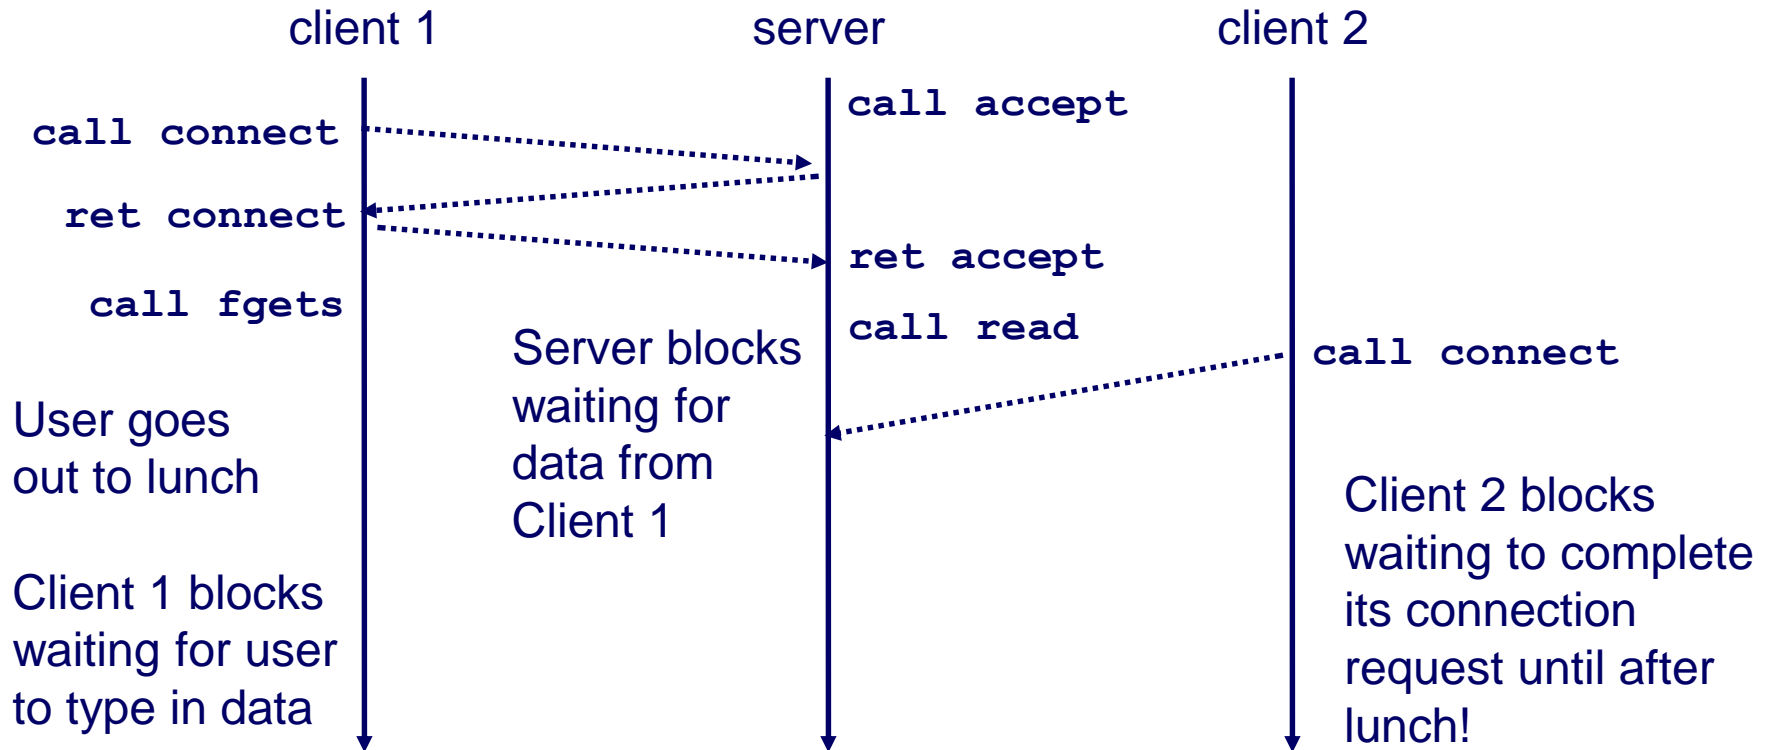
# Revisited



# Close the Socket

- Don't forget to close the socket descriptor, like a file
  - *int close(int sockfd);*
- Now server can loop around and accept a new connection when the old one finishes
- What's wrong here?

# Server Flaw



# Concurrent Servers

